

**CS310 Natural Language Processing - Assignment 3: Recurrent Neural Networks –
Language Modeling and Named Entity Recognition
Total points: 60 + (15 bonus)**

Tasks

- 1) Train a vanilla RNN language model on 《论语》 and evaluate its perplexity.
- 2) Train a bidirectional LSTM model on the CoNLL2003 English named entity recognition task set and evaluate its performance.

Submit

- The modified notebook files `A3_rnn-lm.ipynb` and `A3-ner.ipynb`.
- A `write-up document` in Word/PDF reporting your results in Task 1-3 and Task 2-3.
- For each bonus question you have done, submit a `stand-alone notebook`. For example, `A3_HMM.ipynb`, `A3_CRF.ipynb` etc.

Requirements

Task 1 - LM (30 points):

- 1) (5 points) Data preprocessing.
- 2) (15 points) Model implementation.
 - a) Use `torch.nn.RNN` module
 - b) Do NOT use bidirectional network; multi-layer is fine.
- 3) (10 points) Evaluation and extended experiment.
 - a) Report perplexity on *training set*, as the dataset is very small.
 - b) Generate some sentences; quality is not graded.
 - c) Compare the perplexity on two conditions: randomly initialized embeddings vs. with pretrained embeddings (from A2).

Task 2 - NER (30 points + 15 bonus points):

- 1) (10 points) Data preprocessing.
 - a) Load the train, dev, and test data; build vocabularies for words and labels (tags); defined a data loader that return batches. **Note** that it is recommended to convert all words to *lower cases*, because that is how words are stored in the pretrained embeddings.
 - b) Load the pretrained embedding data to initialize the embedding layer in model. **Note** that you only need to load those words that have occurred in your vocabulary.
The URL for the pretrained embedding is: <https://nlp.stanford.edu/data/glove.6B.zip>. It includes dimension 50, 100, 200, and 300. 100-d should be sufficient.
- 2) (10 points) Implement the “level 1” sequential classifier model, with bi-LSTM architecture.
 - a) Use `torch.nn.LSTM` module.
 - b) Adjust hyperparameters such as hidden size, layer numbers etc. as you like. **Note:** bi-directional and multi-layer network is highly recommended.
- 3) (10 points) Train, evaluate, and save.
 - a) Use greedy search to obtain the predicted labels on test set, i.e., pick the highest probability label for each time step.
 - b) Report F-1 score on test set. **Note** that $F1 = 2 \cdot \frac{precision \cdot recall}{precision + recall}$

c) It is not required, but you can add L2 regularization term to increase test performance.

**** Grading rubrics ****

- If your model is implemented correctly and the training code can run without problem, then you get the full credits for step 1) and 2)
- If you achieve 70% (and above) F-1 score on the **test** set, you get full credits for step 3).
- If your F-1 score $x > 0.5$, then you receive $\left\lfloor \frac{x}{0.7} \cdot 10 \right\rfloor + 0.5$ points for step 3).
- If your F-1 score $x < 0.5$, then you receive 0 points for step 3.

**** Grading rubrics for bonus tasks ****

The three bonus tasks are independent of each other.

Some useful resources and existing implementations you can learn from:

- A good repo implementing and explaining biLSTM+CRF:
<https://github.com/sgrvinod/a-PyTorch-Tutorial-to-Sequence-Labeling>
- A beam search implementation in context of PyTorch and seq2seq:
<https://github.com/budzianowski/PyTorch-Beam-Search-Decoding>
- The implementation of bi-LSTM+CRF from PyTorch official tutorial:
https://pytorch.org/tutorials/beginner/nlp/advanced_tutorial.html

DO NOT directly copy their code!

4) **(3 bonus points)** Implement the maximum entropy Markov model (MEMM).

- Create an embedding layer for all the labels (tags).

5) **(4 bonus points)** Implement beam search for decoding at testing time.

- Compare its performance (F-1 score) with step 3).

6) **(8 bonus points)** Implement conditional random field with Viterbi algorithm (for training and decoding).

- Compare its performance (F-1 score) with step 3).